

DSFA

Spring 2018

Lecture 20

Error Probabilities

Announcements

Hypothesis Testing

Testing a Hypothesis

Step 1: Select Two Hypotheses

- A test chooses between two views of how data were generated:
Null hypothesis proposes that data were generated at random;
Alternative hypothesis proposes some effect other than chance

Step 2: Choose a Test Statistic

- A value that can be computed from the data

Step 3: Compute What The Null Hypothesis Predicts

- Compute the distribution of the test statistic: what the test statistic might be if the null hypothesis were true.

Step 4: Compare the Prediction to the Observed Data

Hypothesis Testing Logic

Define two mutually exclusive descriptions: either this or that.
One of them can be evaluated using probability (the null hypo.)
You can "reject the null," so then you accept the alternative.
Otherwise: you're still not sure, but null looks plausible.

Step 1:
Select Two Hypotheses

Example: The Two Hypotheses

Gregor Mendel (1822-1884) was an Austrian monk and founder of the modern field of genetics. Among many experiments, he tested the hypothesis that pea plants will bear purple or white flowers at random, in the ratio 3:1.

- **Mendel's model describes the world.** If the distribution of the observed plants is different from the distribution in the model, it's just chance variation.
- **Mendel's model doesn't.**

Alternative

Null

(Demo)

Example: Smoking and Babies

Researchers are interested in whether there is an association between smoking mothers and the health of their babies. For each birth, they record the baby's birth weight and whether the mother smokes or not.

- **Birth weights aren't affected by maternal smoking.**

The birthweight distribution for babies of smokers is same as that of babies of non-smokers.

Null

- **They are affected.**

Alternative

Example: Smoking and Babies

Researchers are interested in whether there is an association between smoking mothers and the health of their babies. For each birth, they record the baby's birth weight and whether the mother smokes or not.

- **Birth weights aren't affected by maternal smoking.** The birthweight distribution for babies of smokers is same as that of babies of non-smokers.
 - **They are lower.** Birthweight of babies of smokers are lower than birthweights of babies of non-smokers.
-

Step 2:
Choose a Test Statistic

Choosing a Test Statistic

Test statistic: The statistic that you have chosen to calculate, to help you decide between the two hypotheses

Goal: If the null hypothesis is false, then you expect that measuring the test statistic will allow you to reject the null

Choosing a T

For Mendel's pea flower color, which test statistic would be reasonable to use?
(Choose all that are OK.)

If the alternative hypothesis is true, will test statistic be *larger* than prediction, *smaller*, or *could be either way*?

1. The proportion of plants with purple flowers.
2. The proportion of plants with white flowers.
3. $\text{abs}(p - 0.75)$, where p is the proportion of plants with purple flowers.
4. The number of different colors in the plants flowers.
5. The total variation distance between the distribution in the observed data, vs the model distribution $(0.75, 0.25)$

Choosing a T

For the baby birth weight, the following would be reasonable tests.
(Choose all that are OK.)

If the alternative hypothesis is true, will test statistic be *larger* than prediction, *smaller*, or *could be either way*?

1. The average birth weight of all the babies.
2. The proportion of babies whose mother smoked.
3. The average birth weight of babies of smokers, minus the average birth weight of babies of non-smokers.
4. The absolute value of the previous difference.

Absolute Values & Alternatives

- Choose a test statistic where alternative hypothesis suggests which direction statistic will go.
 - **Alternative: Smoking causes poor health.**
 - Test statistic: Average birth weight for smokers, minus average for non-smokers.
 - **Alternative: Smoking has some relation to health.**
 - Test statistic: Absolute value of that difference.
-

Step 3:
**Compute the Distribution of
the Test Statistic under the
Null Hypothesis**

Step 4:
**Compare the Prediction to the
Observed Data**

Conclusion of a Test

Resolve choice between null and alternative hypotheses

- Compare observed test statistic to its empirical distribution under the null hypothesis
- If the observed value is **consistent** with the distribution, then the test *does not* reject the null hypothesis

Whether a value is consistent with a distribution:

- A visualization may be sufficient
 - Convention: The observed significance level (P-value)
-

Quantifying Conclusions

Step 0: Go find some data. These are the *observations*.

Step 1: Two descriptions of the world:

- Null: Data come from a well-defined random process
- Alternative: Something else is going on

We evaluate how unusual the data would be under the null

Step 2: Choose a test statistic to summarize the data.

Step 3: Compute the following probability (p-value)

$P(\text{the test statistic would be equal to or more extreme than the observed test statistic under the null hypothesis})$


Definition of P -value

The P -value is the chance,

- under the null hypothesis,
- that the test statistic
- is equal to the value that was observed in the data or is even further in the direction of the alternative.

Quantifying Conclusions

P(the **test statistic** would be **equal to or more extreme** than the **observed test statistic** under the null hypothesis)



Evaluating Mendel's
pea flower hypothesis

This area is the P-value
(approximately)

Conventions of Consistency

- **“Inconsistent”**: The test statistic is in the tail of the null distribution.
 - **“In the tail,” first convention**:
 - The area in the tail is less than 5%.
 - The result is “statistically significant.”
 - **“In the tail,” second convention**:
 - The area in the tail is less than 1%.
 - The result is “highly statistically significant.”
-

Sir Ronald Fisher, 1890-1962



"We have the duty of formulating, of summarizing, and of communicating our conclusions, in intelligible form, in recognition of the right of other free minds to utilize them in making their own decisions."

Ronald Fisher

Sir Ronald Fisher, 1925

“It is convenient to take this point [5%] as a limit in judging whether a deviation is to be considered significant or not.”

— *Statistical Methods for Research Workers*

Sir Ronald Fisher, 1926

“If one in twenty does not seem high enough odds, we may, if we prefer it, draw the line at one in fifty (the 2 percent point), or one in a hundred (the 1 percent point). Personally, the author prefers to set a low standard of significance at the 5 percent point ...”

Sir Ronald Fisher, 1935

“No isolated experiment, however significant in itself, can suffice for the experimental demonstration of any natural phenomenon.”





P-hacking

Demo: <https://projects.fivethirtyeight.com/p-hacking/>

Solution: replicate the experiment

Can the Conclusion be Wrong?

Yes.

	Null is true	Alternative is true
Test rejects the null		
Test doesn't reject the null		

(Demo)

An Error Probability

- The cutoff for the P-value is an error probability.
 - If:
 - your **cutoff is 5%**
 - and the **null hypothesis happens to be true**
 - (but you don't know that)
 - then there is about a **5% chance** that **your test will reject the null hypothesis anyway**.
-

Assess this:

“Statistical significance is an objective, unambiguous, universally accepted standard of scientific proof.

— Letter to *Nature*, 1994
